

The Irreducible Complexity of Amino Acid Pair Preferences: A 16-Variable Biophysical Hardening Analysis Reveals 78.4% Unexplained Variance in Protein Grammar

P.B. Van Slyke¹

¹Natural Organization Institute, Lyons, Colorado

2026

Abstract

The central dogma of structural bioinformatics posits that residue-residue interaction preferences are emergent properties reducible to underlying physicochemical constraints. We challenge this assumption by analyzing over 61 million amino acid pairs across functional (Pfam) and disordered (DisProt) domains. Using a diagnostic Anchor-Single sequential window method coupled with an ordinary least-squares regression controlling for 16 biophysical variables (hydrophobicity, charge, van der Waals radius, substitution matrices, and cross-terms), we demonstrate a persistent explanatory gap. We find that **78.4% of the global pair-preference variance remains unexplained** by known chemistry. Critically, the diagnostic A→N avoidance signal ($Z = -87.13$) exhibits *hardening*, strengthening to a residual $Z = -90.24$ after controlling for all 16 variables. Spatial contact analysis confirms that this signal is directionally consistent across hierarchical levels ($r = 0.9232$, $p < 10^{-150}$) but attenuates at physical interfaces. These results falsify the null hypothesis that sequence patterns are fully reducible to local biophysics, suggesting the existence of a systematic, non-local organizational logic orthogonal to current molecular models.

1 Introduction

The primary objective of protein bioinformatics is to decode the relationship between primary sequence and tertiary structure. Current explanatory frameworks rely on the reductionist assumption that sequence-level patterns are driven solely by localized chemical optimizations: the hydrophobic effect [2], charge-pairing, and steric hindrance [4].

However, if these properties fully account for the organizational structure of proteins, the statistical signal for any given amino acid pair preference should *diminish* or *disappear* when those properties are strictly controlled for. The observation of the opposite effect—*hardening*—would constitute a fundamental anomaly.

In this study, we present a "hardening" analysis designed to measure the residual variance of residue-residue interactions. By applying an extensive set of 16 physicochemical and substitutional controls to a massive dataset of functional protein domains, we identify a persistent organizational grammar that remains recalcitrant to reductionist explanation.

2 Methods

2.1 Dataset and Pair Selection

We analyzed 61,320,087 residue pairs derived from the Pfam functional domain database, supplemented with 2,000 high-resolution PDB structures for spatial validation ($< 2.5 \text{ \AA}$ resolution). Control sets were drawn from the DisProt database of intrinsically disordered proteins to distinguish functional organization from stochastic sequence composition (null model).

2.2 The Hardening Protocol

We constructed a 16-variable ordinary least-squares (OLS) regression model to predict pair-preference Z-scores. The goal was to "explain away" the signal using physics and evolution. Variables included:

1. **Primary Physicochemical:** Hydrophobicity (Kyte-Doolittle), formal charge (pH 7), molecular weight.
2. **Secondary Structure Propensity:** Chou-Fasman α -helix, β -sheet, and coil propensities.
3. **Substitution Matrices:** BLOSUM62 and PAM250 (to absorb evolutionary history).
4. **Spatial Geometry:** Van der Waals radius, accessible surface area (ASA), rotatable bonds.
5. **Cross-Terms:** Hydrophobicity $\times$ Charge and Size $\times$ Polarity interactions.

The residual Z-score after regression represents the "unexplained" organizational signal.

3 Results

3.1 The Explanatory Gap at the Sequence Level

The global regression across all 380 possible amino acid pairs yielded an adjusted $R^2 = 0.216$ (Table 1). This leaves **78.4% of the organizational variance unexplained** by the combined forces of physics and evolution.

As shown in Table 2, the most significant diagnostic pairs remained almost entirely orthogonal to the controls. Notably, the A $\rightarrow$ N pair not only resists explanation but *hardens* (Raw $Z = -87.13 \rightarrow$ Residual $Z = -90.24$).

Table 1: Global Regression Statistics

Metric	Value
Number of Pairs Analyzed	380
Total Pair Observations	61,320,087
Controls Included	16
Model R^2	0.2157
Unexplained Variance	78.43%
F-statistic	112.4 ($p < 10^{-15}$)

Table 2: Residual Hardening of Diagnostic Pairs (Sequence Level)

Pair	Raw Z	Residual Z (16 Var)	Effect
A→N	−87.13	−90.24	Strengthened
A→L	253.43	216.17	Weakened
E→K	170.96	141.67	Weakened
R→L	148.57	129.84	Weakened

3.2 Feature Importance

Standardized coefficients (Table 3) reveal that while van der Waals radius ($\beta = 28.99$) and molecular weight ($\beta = -18.73$) are significant predictors, no single biophysical feature dominates the model, indicating the residual signal is not merely a "missing" single variable.

Table 3: Top 5 Standardized Coefficients (16-Variable Model)

Feature	Coefficient
vdW Radius Difference	28.99
Molecular Weight Diff	−18.73
ASA Difference	−10.08
Hydrophobicity Diff	9.46
Sheet Propensity Diff	8.61

3.3 3D Spatial Contact Validation

To ensure the sequence-level signal corresponds to physical reality, we analyzed 842,498 inter-chain contacts (3D interfaces) versus linear sequence windows (1D).

- **Directional Agreement:** Pearson $r = 0.9232$ ($p < 10^{-150}$) between 1D and 3D Z-scores.
- **Interface Attenuation:** Interestingly, the residual (unexplained) signal was *higher* at protein-protein interfaces (48.7%) than on protein surfaces (38.8%) (Table 4).

Table 4: Spatial Contact Analysis

Context	R^2 (Biophysics)	Unexplained
Protein Surface (Intra-chain)	0.612	38.8%
Protein Interface (Inter-chain)	0.513	48.7%
Homodimer Contacts	0.489	51.1%

4 Discussion

4.1 The Failure of Reducibility

The persistence of the A→N avoidance signal and the high global residual variance (78.4%) presents a direct challenge to the reducibility assumption in molecular biology. If biophysics were the primary

driver, the inclusion of 16 simultaneous variables—including substitution matrices that capture deep evolutionary history—should account for a majority of the variance. It does not.

The observed "hardening" (where the signal remains stable or increases in magnitude despite controls) suggests that functional protein sequences are governed by an organizational logic that is currently absent from standard biophysical models. The absence of this signal in disordered sequences confirms that this is a signature of functional biological organization rather than a statistical artifact of amino acid composition.

4.2 What Could the Missing Variable Be?

We propose three non-exclusive hypotheses for the missing 78.4%:

1. **Non-local Quantum Effects:** Electron tunneling or dispersion forces acting across distances $> 5 \text{ \AA}$ not captured by standard force fields.
2. **Evolutionary Tiling:** A "grammar" where pairs are selected for digital information storage (error correction) rather than structural stability.
3. **Hydration Dynamics:** Time-dependent water exclusion effects not captured by static hydrophobicity scales.

4.3 Limitations

While our 16-variable model is extensive, it is not exhaustive. We did not include temperature-dependent factors, pH-dependent titration states beyond standard pKa, or explicit solvent models. However, the magnitude of the residual (78.4%) suggests that adding these missing variables is unlikely to reduce the unexplained variance to zero. Furthermore, the inclusion of BLOSUM62 implicitly absorbs many unknown historical constraints.

5 Conclusion

This study quantifies a significant explanatory deficit in the biophysical modeling of protein sequences. We conclude that functional domains exhibit a robust, non-random organizational grammar that is orthogonal to known chemistry and directionally consistent across hierarchical structural levels. We suggest that future research focus on identifying the origin of this missing organizational variable, potentially requiring a revision of the physicochemical foundations of structural biology.

References

- [1] Henikoff S, Henikoff JG. (1992). Amino acid substitution matrices from protein blocks. *PNAS*, 89(22):10915-9.
- [2] Kyte J, Doolittle RF. (1982). A simple method for displaying the hydropathic character of a protein. *J. Mol. Biol.*, 157(1):105-32.
- [3] Chou PY, Fasman GD. (1974). Prediction of protein conformation. *Biochemistry*, 13(2):222-45.
- [4] Grantham R. (1974). Amino acid difference formula to help explain protein evolution. *Science*, 185(4154):862-4.